

title: "Writing Development as a Provenance-Preserving Scientific Process - Concept Draft"

author: "Dr. William James Stanley"

ai_agent: "GPT-5.6 Sol"

date: "2026-09-11T07:34:10+10:00"

timezone: "Australia/Melbourne"

copyright: "Copyright Dr. William Stanley"

human_gate: "Author retains final authority over meaning, publication, and disclosure."

source_sha256: "see companion .sha256 manifest and PDF footer"

version: "0.1.0"

status: "CONCEPT_DRAFT_INCOMPLETE"

document_type: "method_concept_draft"

scope: "Emerging process for author-AI writing development, provenance, revision, branching, and publication"

epistemic_status: "Preliminary abstraction from one live writing-development sequence; not validated as a general method."

source_boundary: "Current conversation only"

Writing Development as a Provenance-Preserving Scientific Process - concept draft

This document records an emerging writing-development method discovered through live collaboration. It is intentionally incomplete. It is **not** a finished methodology, framework, paper, product, or public specification. Its purpose is to preserve the development process so the author can think further before formalising anything.

Current observation

The writing process that emerged here resembles ordinary scientific-paper development more than one-shot AI drafting. The author begins by speaking or writing ideas in an exploratory form, including contradictions, emotional reactions, technical fragments, analogies, and partial claims. The AI does not treat this first pass as a defective final draft. It treats it as source material whose provenance should be preserved.

A first synthesis then reorganises the material into a coherent argument without claiming that the synthesis is the author's final position. The author takes that synthesis away, reacts to it, identifies what has been flattened or misunderstood, and returns with amendments. Those amendments become new evidence about the intended argument. The next synthesis incorporates them while retaining a visible lineage back to the earlier states.

This produces a development chain rather than a replacement chain: **idea capture -> synthesis -> author criticism/amendment -> reintegration -> publication branching**. Earlier versions remain useful because they show where concepts entered, which phrasing originated from whom, what was rejected, and how the position became more precise.

Layers already emerging

One layer concerns **voice and cadence**. The preferred style is intellectually rigorous but recognisably the author's: direct, occasionally blunt, philosophically inclined, and capable of moving between technical and colloquial registers without turning into generic professional prose. A standing structural rule has also emerged: cohesive paragraphs rather than single-line paragraphing as a stylistic device.

A second layer concerns **vocabulary and disciplinary signalling**. When the subject warrants it, the prose should use the actual vocabulary of the field-mathematics, epistemology, philosophy of science, provenance, benchmark design, scientific norms-because precision matters and because competent terminology signals that the writer understands the terrain. The constraint is that terminology should carry conceptual load rather than become decorative jargon.

A third layer concerns **argument development**. The AI's useful role is not to replace authorship but to expose possible structure: identify the thesis, separate claims from evidence, distinguish philosophical from empirical assertions, reorganise dependencies, and make contradictions visible. The author then accepts, rejects, redirects, or extends that structure. This preserves the Human Gate over meaning and conclusion.

A fourth layer concerns **publication branching**. One developed argument may legitimately produce different descendants. A LinkedIn post can preserve the central controversy and intellectual hook while omitting the full epistemic architecture. A website essay can retain historical context, philosophical development, methodological distinctions, and unresolved questions. These are not merely shortened and lengthened copies; they are different representations with a documented common ancestor.

A fifth layer concerns **IP boundaries**. The public text may explain why the author is using a problem, what epistemic role it serves, and what philosophical principle is being tested without exposing implementation details, private system architecture, procedural mechanics, or unpublished research strategy. Description of purpose is not disclosure of mechanism.

Provenance principle

Every material stage should be recoverable. At minimum, a development record should preserve date, author, AI collaborator, version/state, source lineage, epistemic status, Human Gate status, and cryptographic linkage between working source and rendered artifact. Exact wording should be retained where available; summaries should be labelled as summaries. Unknown timestamps or missing evidence should remain unknown rather than being reconstructed by guesswork.

GitHub/public research-development direction - parked

The author has proposed eventually expressing this process in a GitHub-compatible form resembling scientific paper development: versioned source, visible revisions, provenance, branchable publication outputs, and recoverable history. That idea is **parked**. No repository structure, schema, naming convention, automation, or public specification should be treated as settled yet.

Open questions for the author

- What is the unit of provenance: sentence, paragraph, claim, draft, commit, or some combination?
- How should authorship be represented when the author supplies the idea and the AI supplies wording or structure?
- What constitutes a material revision versus stylistic editing?
- How should rejected language be preserved without overwhelming the readable history?
- Should publication branches remain linked to a canonical argument source, or should each become independently versioned after branching?
- What metadata is essential, and what metadata becomes bureaucratic noise?
- How should spoken/voice-typed idea capture be preserved when transcription itself may contain errors?
- What would make this process scientifically auditable without turning writing into software engineering theatre?

Status: CONCEPT DRAFT - INCOMPLETE. Further author thought is required before formalisation.